

Haris

Controlled AI case preparation

One assigned customer. Independent human approval.

Sulaiman Almkhaizeem

Cybersecurity, Digital Trust & Resilience

A specific AI-security problem

Untrusted notes can reach a model with too much authority.

The proposed workflow is a relationship manager preparing an annual customer-review case for compliance review.

Why now

CORF explicitly addresses AI attacks, inference-data protection and auditability.

Local bank reports show AI adoption. Workflow demand remains unvalidated.

The gateway controls data and actions

Server identity determines the customer.

Known identifiers are removed before local inference.

The model returns text and has no execution authority.

CPU-only local Qwen
Dell XPS 9320 · i7-1280P


Actual prototype screenshot. Synthetic identifiers only.

Allowed ✓		
Decision reason		
Within assigned case scope		
Redactions	Added gateway time	Model time
3	10.373 ms	494.528 ms
What reaches the model		
Summarize the review for [NAME], Civil ID [CIVIL_ID], email [EMAIL].		
Model draft		
Annual review is due, and updated business activity evidence has been requested.		
Evidence hash		
f29e676c39d3644c9caa3bebf9b0cfe9046cbeb365079354d523448c5f30e3d		

Filing requires an independent reviewer

The reviewer sees the exact draft and content digest.

The backend checks scope, role, expiry and digest before committing the case.



One approval, one case.

Measured limits, visible failures

90.9%

Attack block rate

20 of 22 attacks

5%

False-positive rate

1 of 20 benign requests

0

Literal fixture leaks

Cross-customer checks

240/240 decision events. Added gateway p95: 13.193 ms and 13.321 ms.

Two missed attacks and one blocked training quotation remain in the results.

A workflow proposition in a crowded field

Existing components

LiteLLM: gateway and access controls
Presidio: PII identification and redaction
NeMo: runtime guardrails and evidence

Haris must earn its place

Assigned-customer permissions
Approval tied to exact case content
Arabic and English decision evidence

No detector superiority or bank demand is claimed. Integration effort and evidence reconstruction are the hypotheses to test.

Arabic evidence with explicit trust boundaries

Bilingual reasons

A blocked instruction receives a clear Arabic reason and an evidence hash.

An administrator with both key and database can forge local evidence. External anchoring remains necessary.

محظور ×		
سبب القرار		
اكتشفت محاولة لتجاوز التعليمات		
عناصر محجوبة	زمن البوابة الإضافي	زمن النموذج
0	8.256 ms	0 ms
ما يصل إلى النموذج		
لم تُرسل أي بيانات إلى النموذج.		
مسودة النموذج		
لم يُعرض أي مخرج للنموذج.		
بصمة الدليل		
a51e0eb2bc54cadda38e209e6434013a8f6084f98f47afa2c58f8fdd7d5eaa48		

A solo-founder development path

Built

Working API and
bilingual UI
Local model integration
Security tests and
evidence

Next validation

RM and reviewer workflow
Existing bank alternatives
Integration and review
effort

Before deployment

Bank identity and
operations
Independent assurance
Applicable regulatory
review

Queen Mary EEE graduate. Cyber Security MSc at City St George's. First-place team in the completed ROWAD program.

The SIAP ask

Expert feedback on one practical security boundary.

Validate the workflow. Strengthen Arabic and adversarial testing. Define a responsible synthetic trial.

No CBK approval, licence, funding, sandbox entry or adoption is assumed.