

Haris: controlled AI case preparation

Applicant: Sulaiman Almkhaizeem · **SIAP theme:** Cybersecurity, Digital Trust & Resilience · **Stage:** Working prototype

Problem. Staff AI workflows can mix customer information, untrusted notes and tool authority. A manipulated or mistaken model response can become a data disclosure or unauthorized action if access and execution depend on the model. Reviewers need evidence of the identity, policy and exact content that authorized an action.

Why Kuwait. CORF includes AI adversarial defence, inference-data protection and auditability controls (chapter 4, 7.1.2, p.269). NBK and Boubyan report AI adoption. Those facts support relevance; they do not prove this workflow's frequency, cost or buyer demand. No bank interviews or pilot commitments are claimed.

Smallest useful solution. An RM prepares one assigned synthetic customer case. Haris enforces tenant/customer/model/tool permissions and minimizes text before a local model creates a draft. Internal filing requires a separate reviewer to approve that precise content. The model cannot execute tools or grant itself access. Arabic/English decision evidence links the event to its policy and content digest.

Proof. The package includes a working local API and bilingual responsive UI, a real local Qwen model route, a deterministic control-test adapter, independent approval and an internal case action. A 48-case regression corpus exposes 20/22 blocked attacks, 1/20 benign false positives, no literal cross-customer fixture leakage and complete decision evidence over 240 requests. Two independent runs reproduce categorical outcomes. These are synthetic, fixed-case measurements, not production guarantees.

Differentiation to validate. Existing gateways and PII tools already provide core components. Haris's proposition is a narrow customer-entitlement and exact-content review workflow with bilingual evidence. Its commercial value depends on demonstrating easier integration and decision reconstruction than an existing bank solution.

Founder and program ask. My Queen Mary EEE degree, current Cyber Security MSc, first-place ROWAD team experience and measured embedded-systems projects support a bounded engineering build. I seek SIAP mentoring and expert feedback to validate the workflow, improve Arabic/security evaluation and define a responsible synthetic trial. No CBK approval, licensing, sandbox admission, funding or adoption is assumed.

Sources: [CBK SIAP 2026 brief](#); [CORF](#); bank and competitor links in [sources.md](#); measurements in [evidence/run-1/metrics.json](#) and [evidence/run-2/metrics.json](#).